

CENG 218

Design and Analysis of Algorithms

Izmir Institute of Technology

Lecture 7: Sorting in linear time

Slides were mostly prepared using the material provided by Prof. Charles E. Leiserson and Prof. Erik Demaine from MIT


How fast can we sort?

- It depends on the model, i.e. what you can do with the elements.
- All the sorting algorithms we have seen so far are *comparison sorts*: only use comparisons to determine the relative order of elements.
E.g., insertion sort (n^2), merge sort ($n \lg n$), quicksort ($n \lg n$).
- Is $O(n \lg n)$ the best we can do with comparison sort model?

Decision tree can help us answer this question.

Decision-tree example

Sort $\langle a_1, a_2, a_3 \rangle$




In general (n elements)

- Each internal node is labeled $i:j$ for comparing a_i and a_j
- The left subtree shows subsequent comparisons if $a_i \leq a_j$
- The right subtree shows subsequent comparisons if $a_i > a_j$

Decision-tree example

Example: Sort $\langle 9, 4, 6 \rangle$



Each leaf contains a permutation $\langle \pi(1), \pi(2), \dots, \pi(n) \rangle$ to indicate that the ordering $a_{\pi(1)} \leq a_{\pi(2)} \leq \dots \leq a_{\pi(n)}$ has been established.


Decision-tree model

A decision tree can model the execution of any comparison sort:

- It serves as graphical representation of algorithms.
- One tree for each input size n (not so generic).
- View the algorithm as splitting whenever it compares two elements.
- The running time of the algorithm = the length of the path taken.
- Worst-case running time = height of tree.

Decision-tree, Insertion Sort, $n=4$

An uncomplete decision tree to model insertion sort:




Lower bound for decision-tree sorting

Theorem: Any decision tree that can sort n elements must have height $\Omega(n \lg n)$.

Proof: The tree must contain $\geq n!$ leaves, since there are $n!$ possible permutations.

A height- h binary tree has $\leq 2^h$ leaves.

Thus, $2^h \geq n!$


$$\begin{aligned} h &\geq \lg(n!) && (\lg \text{ is mono. increasing}) \\ &\geq \lg((n/e)^n) && (\text{Stirling's formula}) \end{aligned}$$

Worst-case
running time

$$\begin{aligned} &= n (\lg n - \lg e) \\ &= \Theta(n \lg n) && (\lg e \text{ is constant}) \end{aligned}$$

$$h \text{ is } \Omega(n \lg n) .$$

Sorting in linear time

Counting sort: No comparisons between elements.

- Input: $A[1 \dots n]$, where $A[j] \in \{1, 2, \dots, k\}$.
- Output: $B[1 \dots n]$, sorted.
- Auxiliary storage: $C[1 \dots k]$.

I.e. Size of C is the maximum number in A .

Counting sort

for $i \leftarrow 1$ **to** k

$C[i] \leftarrow 0$

for $j \leftarrow 1$ **to** n

$C[A[j]] \leftarrow C[A[j]] + 1 \quad \triangleright C[i] = |\{\text{key} = i\}|$

for $i \leftarrow 2$ **to** k




$C[i] \leftarrow C[i] + C[i-1] \quad \triangleright C[i] = |\{\text{key} \leq i\}|$

for $j \leftarrow n$ **downto** 1

$B[C[A[j]]] \leftarrow A[j]$

$C[A[j]] \leftarrow C[A[j]] - 1$

Counting sort example



Loop 1




	1	2	3	4	5
<i>A</i> :	4	1	3	4	3

	1	2	3	4
<i>C</i> :	0	0	0	0

<i>B</i> :					
------------	--	--	--	--	--

for $i \leftarrow 1$ **to** k
 $C[i] \leftarrow 0$




Loop 2



for $j \leftarrow 1$ **to** n

$C[A[j]] \leftarrow C[A[j]] + 1 \quad \triangleright C[i] = |\{\text{key} = i\}|$




Loop 2



for $j \leftarrow 1$ **to** n

$C[A[j]] \leftarrow C[A[j]] + 1 \quad \triangleright C[i] = |\{\text{key} = i\}|$

Loop 2



for $j \leftarrow 1$ **to** n

$C[A[j]] \leftarrow C[A[j]] + 1 \quad \triangleright C[i] = |\{\text{key} = i\}|$

Loop 2

	1	2	3	4	5
<i>A</i> :	4	1	3	4	3




	1	2	3	4
<i>C</i> :	1	0	1	2

<i>B</i> :					
------------	--	--	--	--	--

for $j \leftarrow 1$ **to** n

$C[A[j]] \leftarrow C[A[j]] + 1 \quad \triangleright C[i] = |\{\text{key} = i\}|$

Loop 2



for $j \leftarrow 1$ **to** n

$C[A[j]] \leftarrow C[A[j]] + 1 \quad \triangleright C[i] = |\{\text{key} = i\}|$

Loop 3

	1	2	3	4	5
<i>A</i> :	4	1	3	4	3

<i>B</i> :					
------------	--	--	--	--	--

	1	2	3	4
<i>C</i> :	1	0	2	2

<i>C'</i> :	1	1	2	2
-------------	---	---	---	---

for $i \leftarrow 2$ **to** k

$$C[i] \leftarrow C[i] + C[i-1] \quad \triangleright \quad C[i] = |\{\text{key} \leq i\}|$$

Loop 3

	1	2	3	4	5
<i>A</i> :	4	1	3	4	3

<i>B</i> :					
------------	--	--	--	--	--

	1	2	3	4
<i>C</i> :	1	0	2	2

<i>C'</i> :	1	1	3	2
-------------	---	---	---	---

for $i \leftarrow 2$ **to** k

$C[i] \leftarrow C[i] + C[i-1] \quad \triangleright C[i] = |\{\text{key} \leq i\}|$

Loop 3

	1	2	3	4	5
<i>A</i> :	4	1	3	4	3

<i>B</i> :					
------------	--	--	--	--	--



	1	2	3	4
<i>C</i> :	1	0	2	2

<i>C'</i> :	1	1	3	5
-------------	---	---	---	---

for $i \leftarrow 2$ **to** k

$$C[i] \leftarrow C[i] + C[i-1] \quad \triangleright C[i] = |\{\text{key} \leq i\}|$$

Loop 4





for $j \leftarrow n$ **downto** 1

$B[C[A[j]]] \leftarrow A[j]$



$C[A[j]] \leftarrow C[A[j]] - 1$

Loop 4





for $j \leftarrow n$ **downto** 1
 $B[C[A[j]]] \leftarrow A[j]$
 $C[A[j]] \leftarrow C[A[j]] - 1$

Loop 4





for $j \leftarrow n$ **downto** 1
 $B[C[A[j]]] \leftarrow A[j]$
 $C[A[j]] \leftarrow C[A[j]] - 1$

Loop 4



for $j \leftarrow n$ **downto** 1
 $B[C[A[j]]] \leftarrow A[j]$
 $C[A[j]] \leftarrow C[A[j]] - 1$

Loop 4



for $j \leftarrow n$ **downto** 1

$B[C[A[j]]] \leftarrow A[j]$

$C[A[j]] \leftarrow C[A[j]] - 1$

Analysis

$\Theta(k)$ { **for** $i \leftarrow 1$ **to** k
 $C[i] \leftarrow 0$

$\Theta(n)$ { **for** $j \leftarrow 1$ **to** n
 $C[A[j]] \leftarrow C[A[j]] + 1$

$\Theta(k)$ { **for** $i \leftarrow 2$ **to** k
 $C[i] \leftarrow C[i] + C[i-1]$

$\Theta(n)$ { **for** $j \leftarrow n$ **downto** 1
 $B[C[A[j]]] \leftarrow A[j]$
 $C[A[j]] \leftarrow C[A[j]] - 1$

$\Theta(n + k)$

Running time

- If $k = O(n)$, then counting sort takes $\Theta(n)$ time.
- We have seen that *comparison sorting* takes $\Omega(n \lg n)$ time.

Q. How is this possible?


A. Counting sort is not a *comparison sort*.

In fact, not a single comparison between elements occurs!

Note: Counting sort is not a good choice if $k \gg n$.

Stable sorting

Counting sort is a *stable* sort: It preserves the input order among equal elements.



Exercise: What other sorts have this property?


Radix sort

- *Origin*: Herman Hollerith's card-sorting machine for the 1890 U.S. Census. (See details in Appendix)
- Digit-by-digit sort.
- Hollerith's original(bad) idea: sort the most significant digit first. Causing many bins!
- Good idea: Sort on *least-significant digit first* with an auxiliary *stable* sort.

3	2	9
4	5	7
6	5	7
8	3	9
4	3	6
7	2	0
3	5	5


↗
Counting sort is a good choice

Operation of Radix sort



Correctness of Radix sort

- Assume that the numbers are sorted by their low-order $t - 1$ digits.
- Sort on digit t




Correctness of Radix sort

- Assume that the numbers are sorted by their low-order $t - 1$ digits.
- Sort on digit t
 - Two numbers that differ in digit t are correctly sorted.



Correctness of Radix sort

- Assume that the numbers are sorted by their low-order $t - 1$ digits.
- Sort on digit t
 - Two numbers that differ in digit t are correctly sorted.
 - Two numbers equal in digit t are put in the same order as the input \Rightarrow correct order.
(requires an auxiliary stable sort, it may be counting sort but there are other options as well)




Analysis of Radix sort

- Assume counting sort is our auxiliary stable sort.
- So, sorting each digit takes $\Theta(n)$ since k is small (takes values between 0 and 9)
- Then, we need to pass through n numbers for each digit. (Three passes in the previous example).
- If the number of digits is low, it seems good.
- Let's make a more formal analysis.

Analysis continued

- Sort n computer words of b bits each.
- Each word can be viewed as having b/r base- 2^r digits.

Example: 32-bit word



$r = 8 \Rightarrow b/r = 4$ passes of counting sort on base- 2^8 digits;

or $r = 16 \Rightarrow b/r = 2$ passes of counting sort on base- 2^{16} digits.

How many passes should we make?

Analysis continued

- **Recall:** Counting sort takes $\Theta(n + k)$ time to sort n numbers in the range from 0 to $k - 1$.
- If each b -bit word is broken into r -bit pieces, each pass of counting sort takes $\Theta(n + 2^r)$ time.
- Since there are b/r passes, we have

$$T(n, b) = \Theta\left(\frac{b}{r}(n + 2^r)\right)$$

- Choose r to minimize $T(n, b)$:
Increasing r means fewer passes,
but as $r \gg \lg n$, the time grows exponentially.

Analysis continued

$$T(n, b) = \Theta\left(\frac{b}{r}(n + 2^r)\right)$$

- We could minimize $T(n, b)$ by differentiating and setting to 0. (A calculus approach)
- But it is more intuitive to use the observation that we don't want $r > \lg n$.
- Choosing $r = \lg n$ implies $T(n, b) = \Theta(bn/\lg n)$.
- Numbers are in 0 to 2^b range. Lets substitute 2^b with n^d . (observe $d \ll b$). Then, we have $b = d \lg n \Rightarrow$ radix sort runs in $\Theta(dn)$ time.

Analysis result

- Note that $\Theta(dn)$ is the same with $\Theta(bn/\lg n)$ but it is easier to grasp the complexity with $\Theta(dn)$. (This assumes we select the best r and remember $2^b = n^d$).
- Counting sort handles numbers from 0 to n in linear time. Radix sort handles number from 0 to n^d in linear time, which is better.
- As long as $d < \lg n$, radix sort beats *comparison sorting*.

Example comparison with merge sort

Lets say there n many 32-bit numbers ($b=32$):

Assume that we select $r \sim \lg n$ (an optimum choice)

	<u>Radix Sort</u>	<u>Merge Sort</u> ($n \lg n$)
$n=2000$	$2^{32}=2000^d \quad d \sim 3 \quad \Theta(3n)$	$2000 \lg 2000 \sim \Theta(11n)$
$n=256$	$2^{32}=256^d \quad d=4 \quad \Theta(4n)$	$256 \lg 256 \quad \Theta(8n)$
$n=32$	$2^{32}=32^d \quad d \sim 6 \quad \Theta(6n)$	$32 \lg 32 \quad \Theta(5n)$
$n=2000$ $b=16$	$2^{16}=2000^d \quad d \sim 1.5 \quad \Theta(1.5n)$	$2000 \lg 2000 \sim \Theta(11n)$

Conclusion

In practice, radix sort is fast for large size inputs when numbers are small.

Downside: Unlike quicksort, counting sort is not memory friendly. Radix sort (using counting sort) is usually slower than a well-tuned quicksort in practice.

Chapter 8 of the textbook is on linear time sorting. Please try to solve exercises in 8.1 and 8.2.

Appendix: Punched-card technology

- The 1880 U.S. Census took almost 10 years to process.
- While a lecturer at MIT, Herman Hollerith (1860-1929) prototyped punched-card technology.
- His machines, including a “card sorter,” allowed the 1890 census total to be reported in 6 weeks.
- He founded the Tabulating Machine Company in 1911, which later merged with other companies to form International Business Machines (IBM).

Appendix: Punched cards

- Punched card = data record.
- Algorithm = machine + human operator.
- An operator inserts a card into the press.
- Pins on the press reach through the punched holes to make electrical contact with cups beneath the card.
- Whenever a particular digit value is punched, the lid of the corresponding sorting bin lifts. The operator deposits the card into the bin and closes the lid.
- When all cards have been processed, the front panel is opened, and the cards are collected in order.
- Check <http://www.oz.net/~markhow/writing/holl.htm>
- Punch-card creator <http://www.facade.com/legacy/punchcard/>